

ABINIT School 2019

A newcomer-oriented school to ab initio nanoscience simulations


January 21-25, 2019 - Bruyères-le-Châtel, France

PARALLELIZATION IN ABINIT

SPEED-UP EFFICIENCY MANYCORE
MPI OPENMP

Marc Torrent

CEA, DAM, DIF, France



```
autoparal      1
npband        108
bandpp         32
export OMP_NUM_THREADS=32
mpirun -n 108 abinit
```

OUTLINE

What are super-computers made of?

Parallel computer technologies

Scalar and *manycore* processors

How to measure the parallel efficiency

Speedup, efficiency, scalability

ABINIT parallelization strategy

Time consuming parts

Iterative diagonalization algorithm

Parallelism inside ABINIT

Paralelization levels

Workload distribution

Performance (examples)


**WHAT ARE SUPER-COMPUTERS
MADE OF?**

HOW TO USE THEM?

WHAT IS PARALLEL COMPUTING?

■ Easy to say...

- Simultaneous use of **multiple compute resources** to solve a **computational problem**



■ ... but not so easy to implement!

- The problem has to be solved in multiple parts which can be solved **concurrently**
- Each part is associated to a series of **instructions**; instructions are compute processes or memory transfers
- Instructions of each part are executed simultaneously on different Compute Processing Units

■ Traditional Measure of computing performance: FLOPS

- Floating point OPerations per Second

■ How to increase the FLOPS of a computer?

- Do more operations per second
 - Increase the frequency!
- Do floating point operations simultaneously (overlap)
 - Vectorization!

THE POWER COST OF FREQUENCY

The power cost of frequency

	Cores	Hz	(Flop/s)	W	Flop/s/W
Superscalar	1	1.5 ×	1.5 ×	3.3 ×	0.45
Multicore	2	0.75 ×	1.5 ×	0.8 ×	1.88

- Power increase as Frequency³
 - Clock rate is limited
- Power is a limited factor for supercomputers
 - Around 3-5W per CPU nowadays
- Multiple slower devices are preferable than one superfast device!
 - Multiples computing units per CPU!
- More performance with less power?
 - software problem!


■ What is vector computing?

- Vectorization can be considered as a “hardware parallelization”, directly implemented in the processor unit.
- It generalizes operations on scalars to apply to vectors.
- Operations apply at once to an entire set of values.
- The processor uses a specific set of instructions:
Advanced Vector Extensions (AVX)
- The size of vectors is hardware dependent.
Recent processors use 512 bits vectors


VECTORIZATION

- Common vector operations (Implemented in all Hardware)
 - Addition, Multiplication
 - FMA (Fuse Multiply-Add) $a \leftarrow a + (b \times c)$

- Example for addition



- Size of “vector” in recent hardware is increasing



VECTORIZATION

■ Vectorization improve performance but...

- Needs more transistors per surface unit in the chip
- Power and heat accumulation increase linearly with vector size
- → Frequency needs to be reduced!
- Needs suitable code!

Vectorizable

```
DO II=2,NMAX  
  A(II) = B(II)+C(II)  
  D(II) = E(II)-A(II-1)  
END DO
```

Not vectorizable

```
DO II=2,NMAX  
  D(II) = E(II)-A(II-1)  
  A(II) = B(II)+C(II)  
END DO
```

Need A before it has been computed

- Needs code changes to help the compiler!
 - Beware to data interdependency
- Order of operations is non deterministic
 - Round-off errors are unpredictable

IMPROVE PERFORMANCE OF MEMORY TRANSFER

■ Traditional Measure of memory performance:

- Access time (ps)
- Transfer speed, bandwidth (Byte/s)
- Latency (ns)

■ How to speed up memory access?

- Speed up the access
 - Change the technology
 - Implant the memory closer to processing unit!
- Increase bandwidth
 - Increase memory clock rate, increase number of “channels”!
- Decrease the latency
 - Improve “switches”, improve network speed

THE POWER COST OF MEMORY TRANSFER

■ Some facts about memory transfer:

- Bandwidths : CPU cache: 40 GB/s, RAM: 20 GB/s, network: 3.5 GB/s
- Memory evolves less than computational power:
90's: 0.25 FLOPS/Byte transferred, nowadays: 100/Byte transferred

■ Cost of data movement

- Computation of a FMA costs 50 pJ
- Move data in RAM costs 2 nJ
- Communicating data (network) costs 3 nJ


■ Random vs strided access

- Random access is very low ~ equivalent to 200 CPU cycles
- Strided access triggers prefetchers, reduces the latency

■ Recomputing data is faster than fetching it randomly in memory


HOW TO DISTRIBUTE DATA IN MEMORY

Distributed Memory



- Private memory
- Processors operate independently
- Data transfer should be programmed explicitly (MPI)
- Relies on network performances

Shared Memory



- Memory is common to all processors
- Tasks operate concurrently on data
- Relies on bandwidth performances

WHAT DO WE WANT?


If we could have **N** processing units (compute+memory),
we would like one calculation be be **N times faster!**

2 workers are twice faster than one!

What is a worker on a (super)computer?

A BASIC PROCESSOR DESIGN (OLD FASHIONED)


- Arithmetic and Logic Unit
- Floating-Point Unit
- Memory (small)
- Controllers
- ...
- RAM is far from the processor
- 1 processor (CPU) has 1 core!



BUILDING A SUPERCOMPUTER WITH OLD CPUS

MPI = Message Passing Interface

A communication protocol



Message Passing paradigm


- Distributed memory model: process X cannot access RAM
- For 100,000 CPUS, need for a very efficient communication network!

How to use it?


- Install a MPI library and compile the code with it.
- Launch:
`mpirun -n N executable`

A MODERN PROCESSOR DESIGN: « MANYCORE »

- 1 processor has several cores
(nowadays: 4 to 68)
- 1 core = ALU/FPU/cache memory
- Each core may have 2 threads
(concurrent tasks)
- All the cores share the RAM memory
- Core can be slow but highly vectorizable
- Note : the core may be grouped by “sockets”. Sharing memory is easy inside a socket; it is not from one socket to the other
→ Non Uniform Memory Access (NUMA)



BUILDING A SUPERCOMPUTER WITH MODERN CPUS



Hybrid parallelism


- Distributed memory between nodes
- Shared memory inside a node
(beware to NUMA effect)
- Need to know the computer architecture to run a code!

How to use it?

- Select the number of concurrent tasks on a node (*openMP*):
`export OMP_NUM_THREADS=x`
- Launch the code in hybrid mode:
`mpirun -n N -c x executable`

SUPER-COMPUTERS EFFICIENCY INCREASES CONTINUOUSLY ...

Super-computers
world TOP 500



...CODES MUST BE IN CONTINUOUS EVOLUTION

Exercise:

- Take a given computational problem
- Write a code at a time t_0 .
 Solve the problem on a computer.
- Freeze your code and wait some time $t_1 - t_0$
- Take a **new** computer at time t_1 .
 Solve again the same problem.
- **What happens to your performances?**

HOW TO MEASURE THE EFFICIENCY OF A CODE ON A SUPER-COMPUTER?

3 WAYS TO DEAL WITH PARALLELISM

- Speedup
- Scaling efficiency
- Scalability

- *These performance indicators will tell us how good/efficient the parallelism is.*
- *Is the code adapted to massive parallelism?*
- *Do we correctly use it?*


SPEEDUP

For a given case test

The speedup is defined by

$$S(P) = \frac{T(1)}{T(P)} \quad (1)$$

where $T(P)$ is the execution time on P cores.



- The closer to the straight line, the better
- Hyper speedup : cache/memory effect
- Bad speedup : time consuming communication, not enough parallel parts, ...

SPEEDUP – AMDAHL'S LAW

What if the code is only parallelized at $\alpha\%$?

The sequential execution time is:

$$T = (1 - \alpha)T + \alpha T$$

On P cores the time will be at best :

$$T(P) = \underbrace{(1 - \alpha)T}_{\text{sequential}} + \underbrace{\frac{\alpha}{P}T}_{\text{parallel}}$$


Thus, the speedup will be :

$$S(P) = \frac{T(1)}{T(P)} = \frac{T(1)}{(1 - \alpha)T(1) + \frac{\alpha}{P}T(1)} = \frac{1}{1 - \alpha + \frac{\alpha}{P}}$$

SPEEDUP – AMDAHL'S LAW

Theoretical limit of the speedup


$$S(P) = \frac{1}{1 - \alpha + \frac{\alpha}{P}} \quad (2)$$



SCALING EFFICIENCY

The scaling efficiency η is defined as :

$$\eta = \frac{S(P)}{P} = \frac{T(1)}{PT(P)} \begin{cases} \in [0; 1] & \rightarrow \text{normal} \\ > 1 & \rightarrow \text{hyper} \end{cases} \quad (3)$$



SCALABILITY

What is a scalable code ?

There are 2 ways of defining the scalability :

- Strong scaling : The work load is the same but the number of workers increase :

$$\eta = \frac{S(P)}{P} = \frac{T(1)}{PT(P)} \text{ should stay close to 1.}$$


- Weak scaling : The work is increased in the same way as the number of workers :

$$S(P) = \frac{T(1)}{T(P)} \text{ should stay close to 1}$$

TEST ON MATRIX-MATRIX MULTIPLICATION

Test of the so called GEMM

GEneral Matrix Matrix product




Good scaling is reached only if each “worker” has enough data to work on !

ABINIT PARALLELIZATION STRATEGY


ABINIT – WHERE IS SPENT TIME?

The main self-consistent loop




ABINIT – WHERE IS SPENT TIME?

The main self-consistent loop



ABINIT – WHERE IS SPENT TIME?

The main self-consistent loop



MOST TIME CONSUMING PARTS

- Use *timopt=-3* in input file to obtain detailed summary of time spent in ABINIT
- A typical time analysis:

- **vtorho**: $v \rightarrow \rho(r)$ (98%)
 - **vtowfk**: $v \rightarrow |\psi_k\rangle : \hat{\mathcal{H}}_k|\psi_k\rangle = \varepsilon|\psi_k\rangle$ (97%)
 - ▶ **cgwf, lobpcg**: Diagonalization (7%)
 - ▶ **getghc**: $\hat{\mathcal{H}}|\psi\rangle$ (90%)
 - **mkrho**: $n(r) = \sum_i^N \langle \psi_i(r) | \psi_i(r) \rangle$
- **rhotov**: $n(r) \rightarrow v$


- Parallelism efficiency is dominated by
 - Algorithm to find eigenvectors $\hat{\mathcal{H}}_k|\psi_k\rangle = \varepsilon|\psi_k\rangle$
 - Hamiltonian application $\hat{\mathcal{H}}_k|\psi_k\rangle :$

INTERNAL ALGORITHME SCALABILITY

- Target computations : $1000 < N_{band} < 50\ 000$, $N_{PW} \sim 100\ 000 \dots 250\ 000$
- Direct diagonalization unachievable ($\sim 10^{12}$)
- In search of the eigenvectors associated with the lowest eigenvalues
- Need an **iterative algorithm**
- Different kind of algorithms but the elementary bricks are the same:
 - Hamiltonian application (linear algebra, FFT) \triangleright computation
 - Rayleigh-Ritz procedure (linear algebra, diago, ortho) \triangleright communication

INTERNAL ALGORITHME SCALABILITY

Scalability of ABINIT internal algorithm



1 H application per iteration
1 complete orthogonalisation per iteration

1 H application per iteration
1 partial orthogonalisation per iteration
1 complete orthogonalisation

5 H applications per iteration
1 complete orthogonalisation


20 H applications per iteration
No orthogonalisation

ABINIT STRONG SCALING - 2016

Repartition of time in in a ground-state calculation
varying the number of band CPU cores
(MPI, strong scaling)

TEST CASE

A vacancy in a 108 atoms cell (gold)
Gamma k-point only, PAW
Computation of total energy and forces



Iterative solver (eigenvalues)
Without Hamiltonian application


Hamiltonian application
Linear algebra, FFT

ABINIT ON MANY CORE ARCHITECTURES, STRATEGY

1. Improve the scalability of the diagonalization algorithm
More calculation, less storage, less communications
2. Efficiently use the shared memory (*openMP*)
in a « medium grained » mode
Adapt the code to the hardware topology
Give longer tasks to the elementary computing units
Decrease the data movements
3. Externalize the elementary operations
Express the physics in terms of elementary operations
Use vendor (or optimized) libraries
4. Add an abstraction layer
Isolate low level optimized operations

→ Decrease Time To Solution

ABINIT PERFORMANCES - 2019



Iterative solver

Hamiltonian

Old implementation
MPI only

CEA-TGCC Curie
Nehalem
16 cores/node


A perfect example
of Amdahl's law

CEA Tera1000-2
Intel KNL
64 cores/node

Uranium
128 atoms
1600 bands (3200 electrons)

2018

ABINIT - 24 threads



Iterative
solver

Hamiltonian


New implementation
MPI x 24 threads

FROM DISTRIBUTED TO HYBRID PARALLELISM

Distributed (pure MPI)

2016

TGCC – Curie
Intel Nehalem




Titanium
512 atoms (8200 electrons)
Speedup

Hybrid (MPI+openMP)

2019

TGCC – Joliot-Curie
Intel Skylake – 48 cores/node – 2 sockets



Gallium oxide Ga_2O_3
1960 atoms (17400 electrons)
Time per SCF iteration

PARALLELISM INSIDE ABINIT

ELECTRONIC DENSITY AND PARALLELISM

Electronic density formula within PAW+plane-wave DFT

$$\rho(\vec{r}) = \sum_{\sigma} \sum_{\text{spins}}^n \left[\int_{\text{Reciprocal space}} \left(\sum_{\vec{g}} \left| C_{n,k}(\vec{g}) \cdot e^{i \cdot (\vec{k} + \vec{g}) \cdot \vec{r}} \right|^2 \right) \cdot d\vec{k} \right] + \sum_{\text{atoms}}^A \rho_{\text{compensation}}(\vec{r})$$

Parallelization levels:

Spins

Bands

k vectors

Plane waves

Atoms

System-size scaling:

≤ 2

$\propto S$

$\propto 1/S$

$\propto S$

$\propto S$

Parallel speedup:

Linear

???

Linear

Poor

Linear

- Each \mathbf{k} independent from the others → Try first parallelization over \mathbf{k} points
- In case of polarization (nsppol=2) → parallelization over spins

PARALLELIZATION OF EIGENSOLVER

- Eigenvectors are orthogonal to each other → non trivial parallelization
- Can parallelize over bands and/or plane waves
- 2 kinds of algorithms :
 - Treat blocks of eigenvectors concurrently
 - Filter independently sets of eigenvector

- Possible choices of algorithm:
 - **Conjugate Gradient**
Default when k-points parallelization only
 - **Block conjugate gradient** (LOBPCG)
Default when Band-FFT parallelization
 - **Chebyshev Filtering**
For a very large number of processors

ABINIT INPUT VARIABLES FOR PARALLELISM

- By default, only parallelization over k/spins is activated
 - **paral_kgb:** **1** → To activate k/plane-waves/bands parallelism

- Choice of iterative eigensover
 - **wfoptalg:** **0** → conjugate gradient
114 → block conjugate gradient
1 → Chebyshev filtering

- 4 basic input variables to control how things are distributed


npkpt	npband	bandpp	npfft or tasks (openMP)
# procs for k-points	# procs for bands	bands per process	# procs or # tasks for plane-wave/FFT

$$n_{proc} = \prod np^*$$

PARALLELISM OF BLOCKED ALGORITHM (LOBPCG)

■ Diagonalization per block:


- Each block of eigenvectors is concurrently diagonalized by the **npband** processes. One bloc after the other.
- Each process handles **bandpp** bands
→ The size of a bloc is **bandpp x npband**



PARALLELISM OF BLOCKED ALGORITHM (LOBPCG)


- The **accuracy** of the diagonalization depends on the block size:
 - Eigenvectors are orthogonal inside a block;
Then blocks are orthogonalized
The orthogonalization is better for large blocks
 - → ABINIT converges better with large blocks

- The **speed** of the diagonalization depends on the block size:
 - Blocks are diagonalized one after the other
Parallelization is more efficient for one block
 - → ABINIT converges faster with large blocks



TWO INTERNAL REPRESENTATIONS

- Bands are distributed in memory.
Plane waves (FFT components) are distributed in memory
- To perform a scalar product $\langle \Psi_i | \Psi_j \rangle$
a process needs the Ψ_i and Ψ_j vectors and a few plane-waves
- To perform a Fast Fourier Transform (FFT)
a process need all the components of a vector Ψ_i
- We **change the representation** by applying a transposition
→ Needs a lot of communications



FIRST SCENARIO: PURE MPI

- First, if possible, parallelize over k-points/spins (**npkpt**)
- Then parallelize the diagonalization over **npband** x **npfft** processes
 - Each block diagonalization is made by **npband** processes in parallel
 - Each process handles **bandpp** bands sequentially
 - Scalar products and FFTs are done by **npfft** processes
- Distribute the workload as follows:
 - **npfft** should be >1 , but should be small enough (<10)
 - **npband** sets the # of procs: **nproc** = **npkpt** x **npband** x **npfft**
 - Increase the convergence speed by increasing the block size
→ increase **bandpp** but not to much (sequential part)
 - **nband** has to be a multiple of **npband** x **bandpp**
- In any case, the ideal distribution is system dependent!

2ND SCENARIO: HYBRID MPI+OPENMP - 1

- Launch the calculation using a maximum number of *openMP* tasks,
Ideally the size of a “socket” on the compute node
`export OMP_NUM_THREADS=xx`
- First, if possible, distribute MPI processes over k-points/spins (**npkpt**)
- Then distribute the remaining MPI processes over bands (**npband**)
 - Each block diagonalization is made by **npband** processes in parallel
 - Each process handles **bandpp** bands in parallel using the *openMP* tasks
 - Scalar products and FFTs are done by the *openMP* tasks (**npfft** not used)
 - **nproc = npkpt x npband**

2ND SCENARIO: HYBRID MPI+OPENMP - 2

- Distribute the workload as follows:
 - **npband** x **bandpp** (size of a block) should be maximalized and has to divide the number of bands (**nband**)
Ideally it should be **nband** or **nband/2** or **nband/3**
 - **bandpp** has to be a multiple of the number of *openMP* tasks
 - **nband** has to be a multiple of **npband** x **bandpp**
- In any case, the ideal distribution is system dependent!

WORKLOAD DISTRIBUTION – KEY IDEAS

■ Main rules (for large systems):

- Use *openMP* as soon as possible; maximize the # threads
- Use k-points/spins parallelism first
- Maximize the block size (increase **npband** and **bandpp**)
- Follow the multiplicity rules for **nband**, **nthreads**, **npband**, **bandpp**...

■ Make performance comparisons for every system

■ Use **autoparal** and **max_ncpus** variable as a starting guess:

- **autoparal=1**: ABINIT tries to find automatically a suitable distribution
- **max_ncpus=N**: ABINIT prints all possible distributions with **nproc<=N** and stops

ABINIT+MPI+OPENMP - EXAMPLES

Varying the block size

# MPI proc. = npband	# threads = bandpp	# blocks	CPU (sec.)
36	32	1	188
36	16	2	205
36	8	4	231
36	4	8	360

Test case

64 Pu atoms - 1200 electronic bands

Tera1000 CEA super-computer (Intel KNL)

Changing the MPI – openMP distribution

# MPI proc. = npband	# threads = bandpp	# blocks	CPU (sec.)
570	2	1	1091
285	4	1	918
143	8	1	320
72	16	1	248
36	32	1	188
18	64	1	197

OTHER LEVELS OF PARALLELIZATION

- **npimage** → For **images** : NEB, string method, PIMD
- **npspinor** → spinorial components in case of
spin-orbit coupling
- **nppert** → in a DFPT calculation (response) to parallelize
over **perturbations**
- **nphf** → In case of **Hartree-Fock** calculation
- **paral_atom** → Enable or disable parallelization over **atoms**
(automatically set up)

PERFORMANCES

SCALAR VS MANYCORE ARCHITECTURE

Test case


64 Pu atoms

1200 electronic bands

Tera1000 CEA supercomputer

	<i>Nehalem</i>	<i>Haswell</i>	<i>KNL</i>
Process x threads	500 x 1	500 x 1	50 x 4
Number of nodes	63	16	3
CPU time Old implementation	45 min. 36 iterations	27 min. 36 iterations	N/A
CPU time New implementation	10 min. 8 iterations	5 min. 8 iterations	8 min. 8 iterations

ABINIT – MEDIUM SIZE SYSTEM




Gallium oxide Ga_2O_3
1960 atoms
8700 bands (17400 electrons)
Time per SCF iteration

Bands + FFT parallelism only
Can be mixed to other parallelism levels

TGCC – Joliot-Curie
Intel Skylake
48 cores/node – 2 sockets

ABINIT - LARGE SYSTEM



Gallium oxide Ga_2O_3
4160 atoms
18400 bands (**36800 electrons**)
Time per SCF iteration

Bands + FFT parallelism only
Can be mixed to other parallelism levels

TGCC –Joliot-Curie
Intel Skylake
48 cores/node – 2 sockets

CONCLUSION

ABINIT IN PARALLEL – KEYS POINTS

- ABINIT parallel efficiency is strongly **system dependent**
- It is highly recommended to use **hybrid parallelism**
- ABINIT cannot be used without a **minimum knowledge** of...
 - The computer architecture (nodes, CPUs/node, ...)
 - The iterative diagonalization algorithm
- **autoparal** keyword can help
But this is only a starting point!
Manual tuning is always better



Commissariat à l'énergie atomique et aux énergies alternatives

Etablissement public à caractère industriel et commercial | RCS Paris B 775 685 019